

Q6: System Scalability

The system employs several strategies to ensure smooth performance even with millions of simultaneous users:

1. Horizontal Scaling Strategies:

a. Microservices Architecture

- Service-specific scaling based on demand patterns (e.g., scaling Response Service during survey campaigns).
- Independent technology stack optimization for each microservice based on its specific requirements.
- Isolated failure domains prevent system-wide outages when a single service fails.
- Independent deployment cycles allowing updates without complete system downtime.
- Resource allocation optimization with smaller services requiring fewer resources than monolithic applications.

b. Stateless Service Design

- All core services (*Authentication*, *User*, *Document*, *Distribution*, and *Response* Services) are designed to be stateless, allowing for seamless horizontal scaling.
- New service instances can be dynamically added as load increases without any service reconfiguration.
- Each service can be scaled independently based on its specific resource demands.

c. Containerized Deployment

- All services are containerized, enabling easy deployment and scaling across the infrastructure.
- Container orchestration platforms (like Kubernetes) can automatically scale services based on CPU, memory usage, or custom metrics.
- Auto-scaling policies can be defined to proactively increase capacity before peak times.

d. Load Balancing

- Load balancers distribute traffic across multiple service instances.
- Sticky sessions are implemented specifically for WebSocket connections to maintain consistency during real-time collaboration while allowing horizontal scaling.
- Health checks ensure traffic is only routed to healthy instances.

2. Database Scaling Strategies

a. MongoDB Sharded Cluster

- Document data is stored in a MongoDB sharded cluster that horizontally scales by distributing data across multiple shards.
- Sharding is based on document IDs, ensuring even distribution.
- The sharded architecture allows for handling large volumes of CRDT operations and document storage.

b. PostgreSQL Scaling

- User data, permissions, and responses utilize PostgreSQL with read replicas to handle high read loads.
- Connection pooling minimizes the overhead of establishing database connections.
- Vertical scaling can be applied for write-heavy operations while read replicas handle query load.

c. Caching Layer with Redis

- Redis caching significantly reduces database load by storing frequently accessed document snapshots.
- Distributed Redis clusters ensure cache performance scales horizontally.
- Time-to-live (TTL) policies prevent cache bloat and memory exhaustion.

3. Real-time Collaboration Scaling

a. WebSocket Connection Management

- The *Real-time Collaboration* Service uses a hybrid approach to limit WebSocket message frequency:
 - Batching changes based on typing pauses (200ms).
 - Regular interval safety net (2 seconds).
 - Character threshold triggers (20 characters).
- This prevents overwhelming the system during high-frequency editing.
- The system uses sticky sessions for WebSocket connections which ensures that a client remains connected to the same instance of the *Real-time Collaboration* Service.

b. CRDT-Based Collaboration

- CRDT operations are inherently scalable as they don't require global locking.
- The system uses distributed locks only for the minimal time needed to store operations.
- Version vectors track causality without requiring a central coordination point.

c. Asynchronous Processing with Kafka

- Kafka decouples CRDT operation broadcast from persistence, allowing the system to handle operation spikes.
- Kafka's partitioning enables parallel processing of operations.
- The Kafka Consumer Service scales based on partition count to handle high volumes.

4. Peak Load Management

a. Rate Limiting and Throttling

- API Gateway implements rate limiting to prevent abuse and ensure fair resource allocation.
- Throttling policies can be applied to non-critical operations during extreme peak loads.

b. Graceful Degradation

- Non-essential features degrade gracefully under extreme load (e.g., real-time analytics may update less frequently).
- Critical operations (like saving form responses) are prioritized over less critical ones.

c. Regional Deployment

- Services can be deployed across multiple geographic regions to distribute load.
- Users are directed to the closest region to minimize latency and balance global traffic.

5. Monitoring and Predictive Scaling

a. Comprehensive Monitoring

- Real-time monitoring of system performance metrics enables rapid identification of bottlenecks.
- Alerts trigger when predefined thresholds are approached.

b. Predictive Auto-scaling

- Historical usage patterns inform predictive scaling policies.
- The system can pre-scale before anticipated peak times (e.g., when large organizations schedule form distributions).

c. Load Testing

- Regular load testing simulates millions of concurrent users to identify scaling limitations before they affect real users.
- Performance bottlenecks are addressed proactively rather than reactively.