

Q7: Real-Time Collaboration

The system employs several strategies to handle real-time collaboration on form creation and editing. It ensures consistency in what users see and experience:

1. CRDT-Based Collaboration Model

To handle real-time collaboration on form creation and editing while ensuring consistency across users, the system implements a Conflict-free Replicated Data Type (CRDT) approach rather than Operational Transformation (OT). This choice provides several key benefits:

- **Conflict-free merging:**
 - CRDTs mathematically guarantee that concurrent edits made by different Clients can be merged without conflicts, regardless of the order they're received.
- **Eventual consistency:**
 - All connected Clients eventually see the same document state after all operations are applied.
- **Decentralized coordination:** No central authority is required to resolve conflicts.

2. Client-Side Implementation

The client maintains a local in-memory copy of the document and implements a hybrid approach for sending summarised CRDT operations:

- **Optimized Event Parsing Hybrid Approach:**
 - When the user pauses typing (200ms delay).
 - At regular intervals as a safety net (every 2 seconds).
 - Upon reaching a character threshold (20 characters).
- **Local-first editing:**
 - Changes are applied to the local document immediately.
 - Users see their own changes without waiting for server confirmation.
 - Provides a responsive experience even with network latency.
- **Idempotency checks:**
 - Each operation includes a unique ID and version vector.
 - Client checks incoming operations against its history to prevent duplicate application.

3. Server-Side Processing

The Real-time Collaboration Service handles distribution of operations:

- **WebSocket connections:**
 - Maintained with sticky sessions through the load balancer.
 - Enables immediate broadcast of operations to all connected clients.
- **Consistency mechanisms**
 - Distributed locks in MongoDB to prevent race conditions during operation storage.
 - Version vectors track causal relationships between operations
 - Version vectors combined with unique operation IDs to ensure that CRDT operations are applied exactly once (idempotency).
 - Operations are applied in causal order on all clients and servers.
- **Resilience features**
 - Server caches the current document state in memory.
 - Operations are persisted to MongoDB before broadcasting.

- Kafka provides a reliable asynchronous channel for operation distribution.

4. Document Versioning

The system maintains document history and consistency through:

- **Periodic snapshots:**
 - Kafka Consumer Service periodically creates document snapshots.
 - Snapshots are cached in Redis and stored in MongoDB.
- **Version retrieval:**
 - Clients can request specific document versions.
 - Enables form distribution of stable versions regardless of ongoing edits.

5. Ensuring Consistent User Experience

To ensure all users have a consistent experience:

- **Convergence guarantees:**
 - CRDT properties ensure all clients eventually reach the same state.
 - Version vectors maintain causal relationships between edits.
- **Failure recovery:**
 - If a server fails, a new server can reconstruct the document state.
 - If a client disconnects, it can receive missed operations upon reconnection.
 - Document snapshots provide efficient recovery points.
- **Sticky Sessions:**
 - System uses sticky sessions for WebSocket connections, which ensures that a client remains connected to the same instance of the Real-time Collaboration Service.
 - This can help to maintain the order of messages and reduce the complexity of managing real-time updates.