

Q9: Failover and Redundancy

To ensure the high availability of the services and handle potential system failures, I would implement a comprehensive failover and redundancy strategy covering:

1. Service Level Redundancy

- **Stateless Microservices:**
 - Multiple instances of each service deployed across availability zones.
 - Services designed to be stateless, allowing any instance to handle requests.
 - Independent scaling and replacement of service instances.
- **Load Balancing:**
 - Load balancers distribute traffic across healthy service instances.
 - Health checks automatically remove unhealthy instances from rotation.
 - Sticky sessions for WebSockets while maintaining failover capability.
- **Service-Specific Recovery:**
 - Authentication Service: Multiple instances with shared JWT verification capabilities.
 - Document Service: State recovery from MongoDB and version vectors.
 - Real-time Collaboration Service: WebSocket reconnection protocols.
 - Kafka Consumer Service: Offset tracking for recovery after restarts.

2. Data Storage Redundancy

- **MongoDB Cluster Redundancy:**
 - Sharded cluster with replica sets for each shard.
 - Automatic primary/secondary failover within replica sets.
 - Data distributed across multiple availability zones.
 - Distributed locks for coordinating critical operations.
- **PostgreSQL High Availability:**
 - Primary/standby configuration with automatic failover.
 - Synchronous replication for critical data (user credentials, permissions).
 - Read replicas for distributing query loads.
 - Regular backups with point-in-time recovery.
- **Redis Cache Redundancy:**
 - Redis cluster with sentinel for automatic failover.
 - Cross-AZ deployment for resilience.
 - Graceful degradation to MongoDB if Redis is unavailable.
 - Data in Redis is treated as ephemeral with recovery from source data.

3. Messaging System Redundancy

- **Kafka Reliability:**
 - Multi-broker Kafka cluster with replication.
 - Topic replication across multiple brokers.
 - Automated leader election for topics.
 - Consumer group rebalancing for Kafka Consumer Service instances.
- **Kafka Failure Recovery Process:**
 - Document Service detects Kafka unavailability.
 - Fallback to direct MongoDB operation retrieval.
 - Operations continue functioning with degraded performance.
 - Automatic recovery and catch-up when Kafka is restored.

4. System Failure Handling

- **Server Failure Recovery:**
 - New servers retrieve the last processed version vector from MongoDB.
 - Pending CRDT operations are fetched based on timestamp.
 - Latest document snapshot is loaded from MongoDB.
 - Operations are applied in causal order using version vectors.
 - Server becomes consistent with the rest of the system.
- **Partial System Failures:**
 - Circuit breakers prevent cascading failures.
 - Retry mechanisms with exponential backoff for transient issues.
 - Fallback to degraded functionality when dependencies are unavailable.
 - Prioritization of critical operations during recovery.
- **Network Partition Handling:**
 - CRDT-based approach tolerates temporary network partitions.
 - Version vectors track causality across partitions.
 - Eventual consistency when partitions heal.

5. Disaster Recovery

- **Multi-Region Strategy:**
 - Option to deploy to multiple geographic regions.
 - Cross-region data replication for critical data.
 - Regional failover for complete region outage.
- **Backup and Restore:**
 - Regular automated backups of PostgreSQL and MongoDB data.
 - Document version history preserved in MongoDB.
 - Point-in-time recovery capability.
 - Regular disaster recovery testing.
- **Data Integrity Protection:**
 - Immutable document version snapshots.
 - CRDT operations stored with idempotency guarantees.
 - Response data integrity protected with transaction guarantees.